# RoboCasa source baselines: fixed PPO update 30

64 paired cases per task · native success · 600 controls · one attempt per policy

| Task | SFT | PPO30 |
|---|---|---|
| OpenDrawer* | 37/64 | 48/64 |
| CloseToasterOvenDoor | 19/64 | 15/64 |
| CoffeeSetupMug | 10/64 | 21/64 |
| PickPlaceCounterToStove | 40/64 | 38/64 |
| TurnOffStove | 7/64 | 4/64 |
| TurnOnSinkFaucet | 24/64 | 28/64 |

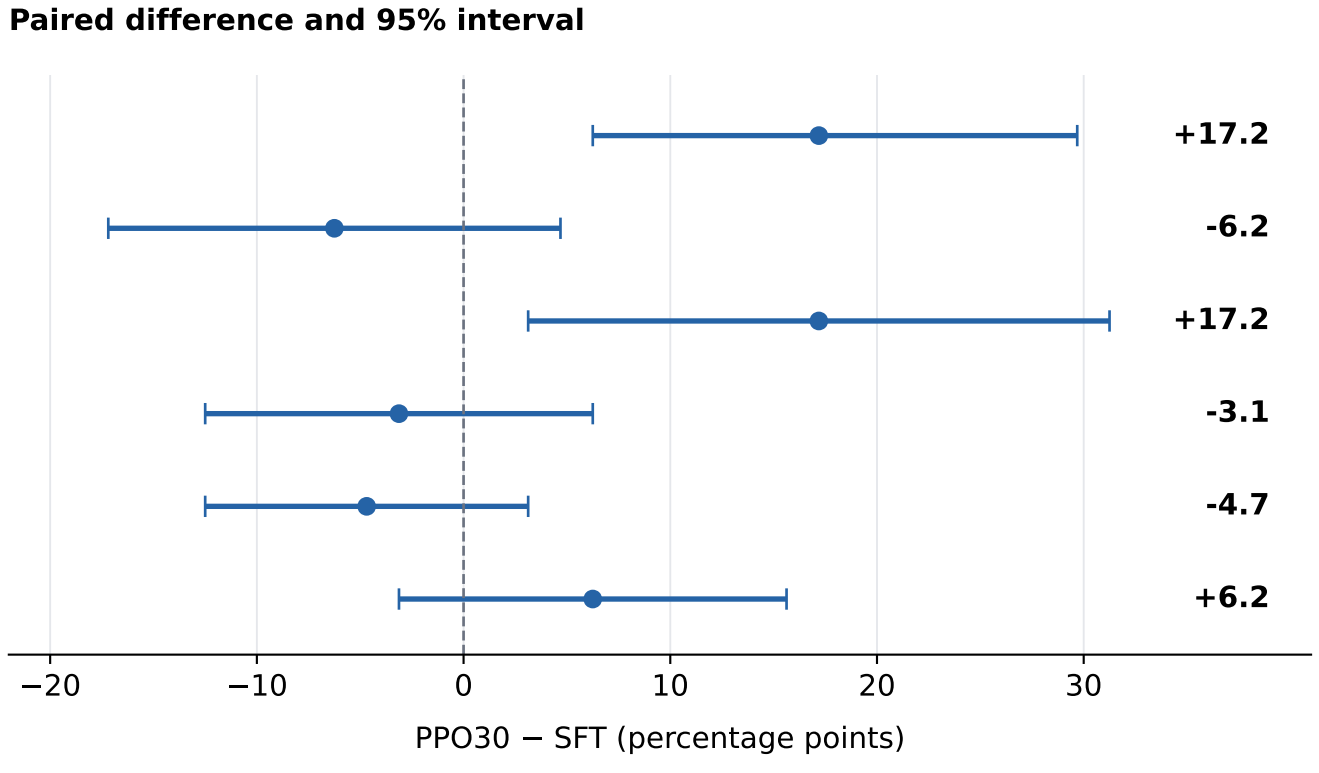


95% paired bootstrap intervals (10,000 draws); marginal, not simultaneous. One trained teacher per task.

*Drawer uses a fresh panel; original SFT 33/64 remains separate. These source baselines do not measure meta-learning or history gains.